STATION 2 · TOKENISE · PRINT ON 2-PART CARBONLESS (NCR) A4, OR PLAIN PAPER OVER CARBON PAPER

TABLE **A** / **B** (circle one — provenance)

TRANSCRIPTION SHEET · DE PAPIEREN MACHINE

EXCERPT No. _____ NAME _____

ONE WORD PER CELL. You decide what counts as a word. A word broken across a line in the original, an abbreviation, a price, a date — you choose, and your choice is the tokenisation. Nobody will correct it.

LEFT TO RIGHT, TOP TO BOTTOM. Read only through the slot in your reading window. Never move it back. Do not fix mistakes: this handwriting is the entire training corpus.

YOU WILL NOT FINISH. Nobody does, and that is designed in. The model uses whatever exists when the timer stops. Leave the rest empty.

THE CARBON COPY MATTERS. Press firmly. The sheet underneath is cut differently from this one, and the two together give **every adjacent pair within each row**.

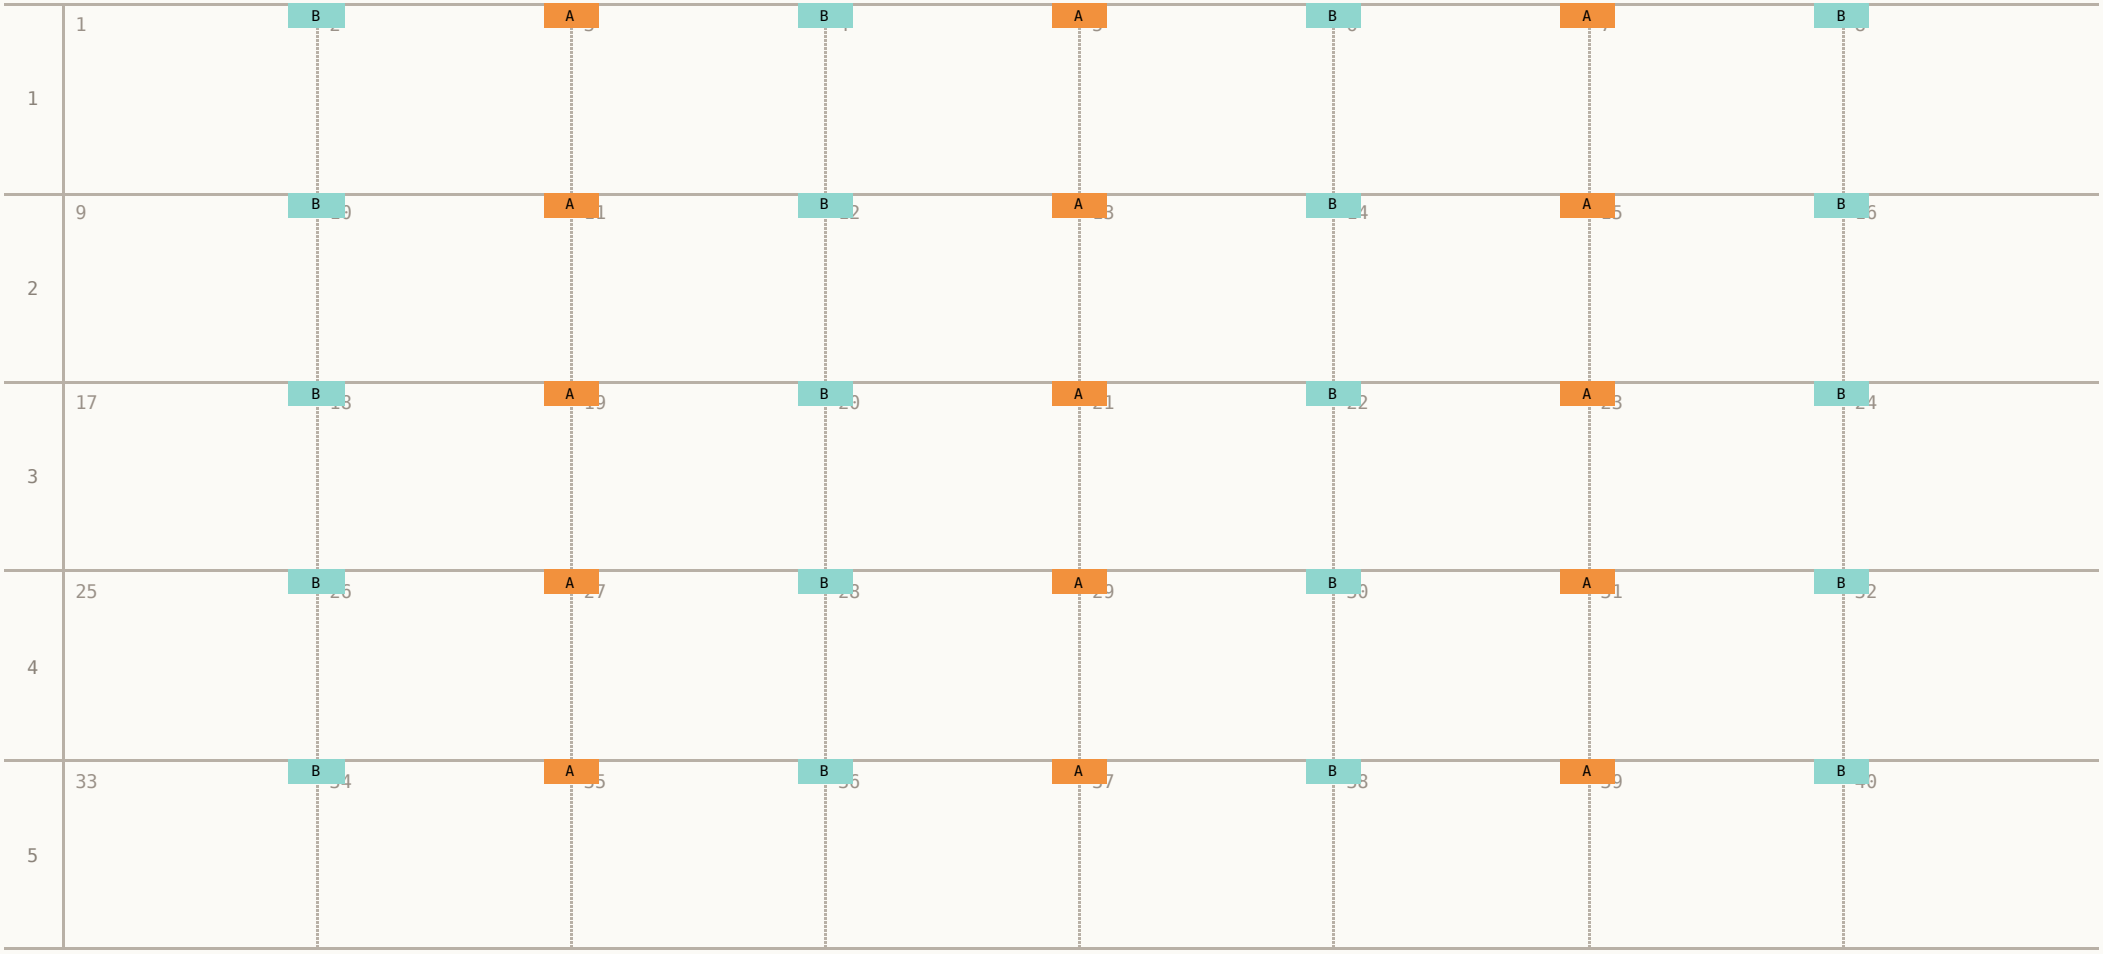

CUTTING (10 MIN):

TOP COPY → cut only at A marks → pairs 1+2, 3+4, 5+6 …

CARBON COPY → cut only at B marks → pairs 2+3, 4+5, 6+7 …

Also cut along every row line.

Single-cell offcuts (10 per sheet) → EXCLUDED box.

WHAT THE CUTTING RULE CANNOT USE. The protocol omits **four cross-row transitions** — 8→9, 16→17, 24→25, 32→33 — and produces **ten single-cell offcuts**, at cells 1 and 8 of each row on the carbon copy. These are two different things. The offcuts are redundant copies of words that do appear in pairs elsewhere; the four transitions are genuinely absent from the model. Both go in the excluded box: they record what the cutting rule could not use. The same artefact breaks text across columns and pages in every digitised archive.

A FULL SHEET → **35 USABLE BIGRAMS** · 4 MISSING TRANSITIONS · 10 OFFCUTS